

# Joke's On You: An Exercise in Joke Generation

CONNOR FORD AND GABE MAGEE

Pomona College

connor.ford@pomona.edu, gabe.magee@pomona.edu

May 3, 2019

## ABSTRACT

Joke generation is a difficult task for humans and machines alike. It combines context, timing, wordplay and world knowledge in a way that can get a reaction out of a real human being. Is there a small predefined subset of a joke that can be generated by machines? We consider a subclass of 'knock-knock' jokes to simplify the generative approach. Using an algorithm as opposed to training and testing more 'intelligently' we are able to create a small number of reasonable jokes with origins from movie scripts.

## I. INTRODUCTION

There have been many successful approaches to joke generation. Cai and Ehrhardt<sup>1</sup> tried to distinguish between a non-joke sentence and a joke one using Neural Nets. Yoshida et al<sup>2</sup> took in various image/caption pairs and tried to produce humorous captions given an image. Finally, Mihalcea and Strapparava<sup>3</sup> tried to apply Linguistic theories about humor to computational generation of one-liner jokes. These jokes typically follow a certain structure like call-and-response, or the more vulgar yo-mama. Others have trained models on large

corpus of data scraped from reddit or twitter. These have less associated structure and generally see more mixed results. We wanted to consider a less common joke-type in current literature: the 'knock-knock' joke. This joke type has a couple main advantages. (i) It follows a rigid structure. To illustrate the format, we annotate the following classic 'knock-knock' joke (not generated).

A: Knock knock.

B: Who's there?

A: Cash. [*Token*]

B: Cash who? [*Token + who = search word*]

A: No thanks, I'll have the peanuts [*Play on search word*]

(ii) There is a discrete set of 'knock-knock' joke subtypes<sup>4</sup>. One of which, as Taylor identifies, is word play on the *token*. We pursue this specific type because it always us to take a straightforward, algorithmic approach to generation. The algorithm looks like:

(a) Generate *Token* from dictionary/corpus by identifying *search words* that end in the 'who'/'ew' and extracting the root of these words.

(b) Use corpus to generate responses based on *search word*.

<sup>1</sup>Cai, J., and Ehrhardt, N. (2013). Is This A Joke?.

<sup>2</sup>Kota Yoshida and Munetaka Minoguchi and Kenichiro Wani and Akio Nakamura and Hirokatsu Kataoka. (2018). Neural Joking Machine: Humorous image captioning.

<sup>3</sup>Mihalcea, R. and Strapparava, C. (2006). Learning to Laugh (Automatically): Computational Models for Humor Recognition.

<sup>4</sup>Taylor, J. (2004). Computationally recognizing word-play in jokes.

## II. METHODS

### i. Generating Tokens

Used corpus of  $\sim 300k$  words and found words ending in the following characters: ["who", "ew", "ewe", "ooo", "ooh", "oo", "hue", "ue", "eau", "eww", "hu"] Then we cross-referenced these words with a pronunciation dataset based on a CMU pronunciation dictionary ending in ["UW0", "UW1", "UW2"]. This only generated a list of 9 potential english words. To extend this, while sacrificing quality, instead of cross-referencing with the pronunciation library, we sacrificed with another english word dictionary. This outputs a list of about 68 english *search words* but the quality of the tokens took a large hit.

### ii. Generating Responses

We decided to generate a response to a joke one of two ways

- (i) Using a bigram learned from a Wikipedia dataset.
- (ii) Using a sentence following approach from a movie script dataset.

## III. RESULTS

token	search word	response
<b>Bigrams</b>		
cash	cashew	cashew apple ii cd of the target for championship at 124.
<b>Sentence Following (1)</b>		
fug	fugue	Fuck that, maybe you're my hallucination.
resid	residue	Putty? On both doors?
mild	mildew	And is that it?
rev	revue	The job's gonna cost you a hundred bucks.
<b>Sentence Following (2)</b>		
wa	wahoo	That's my point!
gl	glue	What's it to you?
curf	curfew	You think he could still be in town?

## IV. DISCUSSION

### i. Why Human Evaluation

We decided to do human evaluation because of the nature of the results. It's really hard to evaluate a joke quantitatively using machine-based methods, due to the lack of models about them. This leaves only a few options, among them chiefly is human evaluation. A score generated from collective human reactions to jokes generated may be a good way to evaluate our models.

### ii. Evaluation Methods

To test our results, we would have a survey that evaluates for different methods for producing jokes. It would include a few generated by each methods along with a place to score the joke as a whole for its coherence and humor.

### iii. Evaluation Results

The survey has not been finished... to be continued.

### REFERENCES

- [Yoshida, Kota et al, 2018] Kota Yoshida and Munetaka Minoguchi and Kenichiro Wani and Akio Nakamura and Hirokatsu Kataoka. (2018). Neural Joking Machine: Humorous image captioning. *CoRR*, abs/1805.11850.
- [Cai, J., and Ehrhardt, N., 2013] Cai, J., and Ehrhardt, N. (2013). Is This A Joke?.
- [Mihalcea, R. and Strapparava, C, 2006] Mihalcea, R. and Strapparava, C. (2006), Learning to Laugh (Automatically): Computational Models for Humor Recognition. *Computational Intelligence*, 22: 126-142. doi:10.1111/j.1467-8640.2006.00278.x
- [Taylor, 2004] Taylor, J. (2006), Computationally recognizing wordplay in jokes *Cognitive Science*