

Joke's On You: An Exercise in Joke Generation

CONNOR FORD AND GABE MAGEE

Pomona College

connor.ford@pomona.edu, gabe.magee@pomona.edu

May 3, 2019

ABSTRACT

Joke generation is a difficult task for humans and machines alike. It combines context, timing, wordplay and world knowledge in a way that can get a reaction out of a real human being. Is there a small predefined subset of a joke that can be generated by machines? We consider a subclass of 'knock-knock' jokes to simplify the generative approach. Using an algorithm as opposed to training and testing more 'intelligently' we are able to create a small number of reasonable jokes with origins from movie scripts.

I. INTRODUCTION

There have been many successful approaches to joke generation. Cai and Ehrhardt¹ tried to distinguish between a non-joke sentence and a joke one using Neural Nets. Yoshida et al² took in various image/caption pairs and tried to produce humorous captions given an image. Finally, Mihalcea and Strapparava³ tried to apply Linguistic theories about humor to computational generation of one-liner jokes. These jokes typically follow a certain structure like call-and-response, or the more vulgar yo-mama. Others have trained models on large

corpus of data scraped from reddit or twitter. These have less associated structure and generally see more mixed results. We wanted to consider a less common joke-type in current literature: the 'knock-knock' joke. This joke type has a couple main advantages. (i) It follows a rigid structure. To illustrate the format, we annotate the following classic 'knock-knock' joke (not generated).

A: Knock knock.

B: Who's there?

A: Cash. [*Token*]

B: Cash who? [*Token + who = search word*]

A: No thanks, I'll have the peanuts [*Play on search word*]

(ii) There is a discrete set of 'knock-knock' joke subtypes⁴. One of which, as Taylor identifies, is word play on the *token*. We pursue this specific type because it always us to take a straightfoward, algorithmic approach to generation. The algorithm looks like:

(a) Generate *Token* from dictionary/corpus by identifying *search words* that end in the 'who'/'ew' and extracting the root of these words.

(b) Use corpus to generate responses based on *search word*.

¹Cai, J., and Ehrhardt, N. (2013). Is This A Joke?.

²Kota Yoshida and Munetaka Minoguchi and Kenichiro Wani and Akio Nakamura and Hirokatsu Kataoka. (2018). Neural Joking Machine: Humorous image captioning.

³Mihalcea, R. and Strapparava, C. (2006). Learning to Laugh (Automatically): Computational Models for Humor Recognition.

⁴Taylor, J. (2004). Computationally recognizing word-play in jokes.

II. METHODS

i. Generating Tokens

Used corpus of $\sim 300k$ words and found words ending in the following characters: ["who", "ew", "ewe", "ooo", "ooh", "oo", "hue", "ue", "eau", "eww", "hu"] Then we cross-referenced these words with a pronunciation dataset based on a CMU pronunciation dictionary ending in ["UW0", "UW1", "UW2"]. This only generated a list of 9 potential english words. To extend this, while sacrificing quality, instead of cross-referencing with the pronunciation library, we sacrificed with another english word dictionary. This outputs a list of about 68 english *search words* but the quality of the tokens took a large hit.

ii. Generating Responses

We decided to generate a response to a joke one of two ways:

- (i) Using a bigram approach learned from a Wikipedia dataset. This meant stringing together similar words to artificially construct a response.
- (ii) Using a 'sentence following' approach from a movie script dataset. This meant finding the occurrence of a word in the movie script data set then returning the next sentence in the dialogue. This dataset is nice because movie scripts allow for more conversational flow then most other corpuses do. It also includes various colloquialisms making it well-suited for jokes.

III. RESULTS

token	search word	response
Bigrams		
cash	cashew	cashew apple ii cd of the target for championship at 124.

Sentence Following (1)		
fug	fugue	Fuck that, maybe you're my hallucination.
resid	residue	Putty? On both doors?
mild	mildew	And is that it?
rev	revue	The job's gonna cost you a hundred bucks.
Sentence Following (2)		
wa	wahoo	That's my point!
gl	glue	What's it to you?
curf	curfew	You think he could still be in town?

IV. DISCUSSION

i. Subsection One

To test our results, we would have a survey that evaluates for different methods for producing jokes. It would include a few generated by each methods along with a place to score the joke as a whole for its coherence and humor.

ii. Subsection Two

The survey has not been finished!

REFERENCES

- [Yoshida, Kota et al, 2018] Kota Yoshida and Munetaka Minoguchi and Kenichiro Wani and Akio Nakamura and Hirokatsu Kataoka. (2018). Neural Joking Machine: Humorous image captioning. *CoRR*, abs/1805.11850.
- [Cai, J., and Ehrhardt, N., 2013] Cai, J., and Ehrhardt, N. (2013). Is This A Joke?.
- [Mihalcea, R. and Strapparava, C, 2006] Mihalcea, R. and Strapparava, C. (2006), Learning to Laugh (Automatically): Computational Models for Humor Recognition. *Computational Intelligence*, 22: 126-142. doi:10.1111/j.1467-8640.2006.00278.x

[Taylor, 2004] Taylor, J. (2006), Computationally recognizing wordplay in jokes *Cognitive Science*